

DAYTONA

Skills vs MCPs

Two layers, one goal: agents that ship.

3 min read


An MCP server hands an agent a box of tools. A Skill teaches it the craft to use them. They are not rivals — they are different layers. Today we are adding one to Daytona: the **Daytona Skill** for Claude.

Different layers, not competitors.

MCP is access. It is the standardized wiring that lets an agent reach an external system. The Daytona MCP gives Claude twelve tools — create a sandbox, run a command, clone a repo, hand back a preview link.

Skills are expertise. A Skill is procedural know-how, loaded only when it is needed. The Daytona Skill teaches Claude *when* to spin up a sandbox, how to chain a build, when to snapshot, and to always clean up afterward.

The difference: tools say what an agent *can* do; a skill says *how* to do it well. You want both.



The cake of agentic capability. Each tier adds a level of expertise — and MCP and Skills are *different tiers*, not competitors.

The Daytona Skill.

It packages everything we have learned about driving secure, sub-90ms sandboxes into a single playbook Claude reads on demand — about 100 tokens of context until the moment it is needed, instead of a dozen tool definitions sitting in every conversation. And it drives the full Daytona SDK: persistent sessions, snapshots, volumes, and code execution — the superset of what the MCP exposes.

MCP gives your agent reach. Skills give it judgment. Together: an agent that runs the code it writes — in a sandbox, and cleans up when it is done.

The Daytona Skill is live. [Documentation](#) · [GitHub](#) · [Sign up](#)